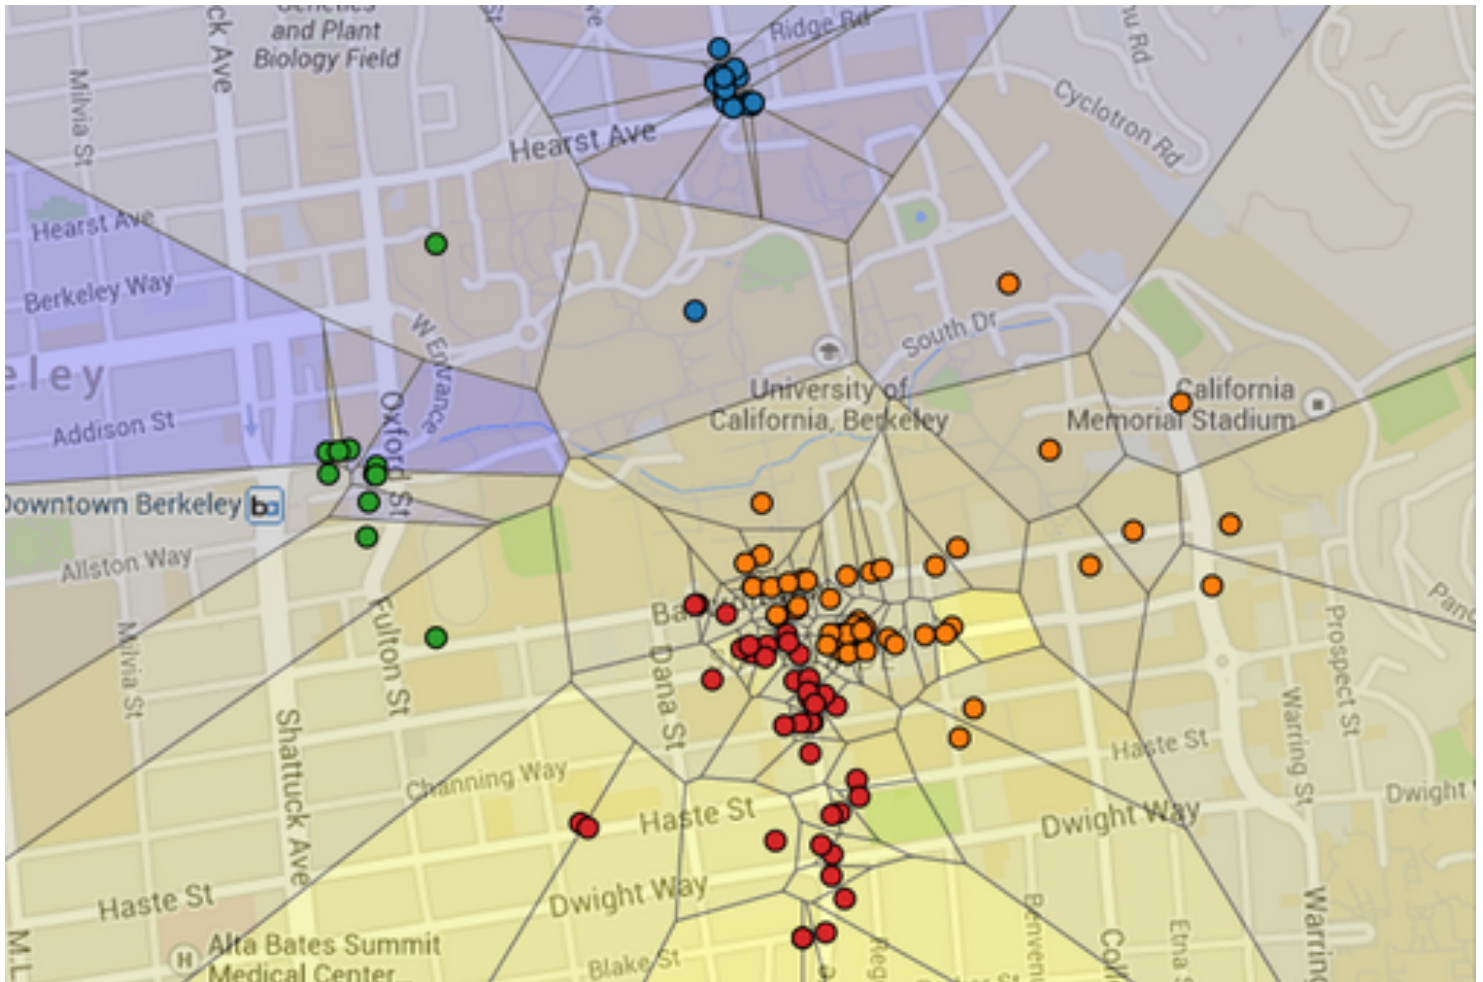


Project 1: Yelp Maps

Yelp Maps

Due by 11:59pm on Sunday, 12/01/19

Download [maps.zip](#) 



Let's go out to eat!

Introduction

In this project, you will create a visualization of restaurant ratings using machine learning and the [Yelp academic dataset](#). In this visualization, Berkeley is segmented into regions, where each region is shaded by the predicted rating of the closest restaurant (yellow is 5 stars, blue is 1 star). Specifically, the visualization you will be constructing is a [Voronoi diagram](#).

In the map above, each dot represents a restaurant. The color of the dot is determined by the restaurant's location. For example, downtown restaurants are colored green. The user that generated this map has a strong preference for Southside restaurants, and so the southern regions are colored yellow.

This project uses concepts from Sections [2.1](#), [2.2](#), [2.3](#), and [2.4.3](#) of [Composing Programs](#). It also introduces techniques and concepts from *machine learning*, a growing field at the intersection of computer science and statistics that analyses data to find patterns and make predictions.

Download starter files

The [maps.zip](#) archive contains all the starter code and data sets. The project uses several files, but all of your changes will be made to `utils.py`, `abstractions.py`, and `recommend.py`.

- `abstractions.py`: Data abstractions used in the project
- `recommend.py`: Machine learning algorithms and data processing
- `utils.py`: Utility functions for data processing
- `ucsd.py`: Utility functions for DSC 20. **DO NOT EDIT**
- `data`: A directory of Yelp users, restaurants, and reviews. **DO NOT EDIT**
- `users`: A directory of user files. **DO NOT EDIT**
- `visualize`: A directory of tools for drawing the final visualization. **DO NOT EDIT**
- `implementations_tester.py`: Utility functions and classes for testing implementation. **DO NOT EDIT TEST FUNCTIONS!**
- `abstractions_tester.py`: Utility functions and classes for testing abstraction. **DO NOT EDIT TEST FUNCTIONS!**
- `test_functions.py`: Helper file for abstractions. **DO NOT EDIT**

Logistics

This is a 2 - week project. You are allowed to complete this project **with one partner**. You should not share your code with any other student (except your partner), or copy from anyone else's solutions.

You will turn in the following files:

- `utils.py`
- `abstractions.py`
- `recommend.py`

For the functions that we ask you to complete, there may be some initial code that we provide. If you would rather not use that code, feel free to delete it and start from scratch. You may also add new function definitions as you see fit.

However, please do **not** modify any other functions. Doing so may result in your code failing our autograder tests. Also, please do not change any function signatures (names, argument order, or number of arguments).

Testing

Throughout this project, you should be testing the correctness of your code. It is good practice to test often, so that it is easy to isolate any problems. However, you should not be testing *too* often, to allow yourself time to think through problems.

There will be two types of testing: implementation and abstraction. It is required that you test both to ensure that you have implemented each method correctly. Each problem will show you how to run the the required tests. All problems test implementation, only some test abstraction.

BE SURE TO PERFORM ALL TESTS PROVIDED.

When you submit, your score is the same as what Gradescope shows you.

You **do not** need to add doctests, docstrings, asserts, etc to your code. The built in testers will be enough.

IMPORTANT: this project is chronological. You cannot move on to future parts

until you complete the previous. You must complete the problems in order.

Phase 1: Utilities

All changes in this phase will be made to `utils.py`.

Problem 1

Before starting the core project, familiarize yourself with some Python features by implementing some functions in `utils.py`.

Each function described below can be implemented in **one line**. These will be helpful in later parts of the project.

Problem 1.1: Using list comprehensions

A list comprehension constructs a new list from an existing sequence by first filtering the given sequence, and then computing an element of the result for each remaining element that is not filtered out.

A list comprehension has the following syntax:

```
[<map expression> for <name> in <sequence expression> if <filter expression>]
```

For example, if we wanted to square every even integer in `range(10)`, we could write:

```
>>> [x * x for x in range(10) if x % 2 == 0] [0, 4, 16, 36, 64]
```

Implement `map_and_filter`, which takes in a sequence `s`, a one-argument function `map_fn`, and a one-argument function `filter_fn`. It returns a new list containing the result of calling `map_fn` on each element of `s` for which `filter_fn` returns a true value. *Make sure your solution is only one line and uses a list comprehension.*

Problem 1.2: Using `min`

The built-in `min` function takes a collection of elements (such as a list or a dictionary) and returns the collection's smallest element. The `min` function also takes in an optional keyword argument called `key`, which must be a one-argument function. The key function is called on each element of the collection, and the return values are used for comparison. For example:

```
>>> min([-1, 0, 1]) # no key argument; return
smallest input -1 >>> min([-1, 0, 1], key = lambda x: x*x) # return
input with the smallest square 0
```

Implement `key_of_min_value`, which takes in a dictionary `d` and returns the key that corresponds to the minimum value in `d`. This behavior differs from just calling `min` on a dictionary, which would return the smallest key. *Make sure your solution is only one line and uses the `min` function.*

Problem 1.3: Using `zip`

The `zip` function defined in `utils.py` takes multiple sequences as arguments and returns a list of lists, where the i -th list contains the i -th element of each original sequence. For example:

```
>>> zip([1, 2, 3], [4, 5, 6]) [[1, 4], [2, 5], [3, 6]] >>> for triple
in zip(['a', 'b', 'c'], [1, 2, 3], ['do', 're', 'mi']): ...
print(triple) ['a', 1, 'do'] ['b', 2, 're'] ['c', 3, 'mi']
```

Use the `zip` function to implement `enumerate`, which pairs the elements of a sequence with their indices, offset by a starting value. `enumerate` takes a sequence `s` and a starting value `start`. It returns a list of pairs, in which the i -th element is $i + \text{start}$ paired with the i -th element of `s`. For example:

```
>>> enumerate(['maps', 21, 47], start=1)

>>> [[1, 'maps'], [2, 21], [3, 47]]
```

Make sure your solution is only one line and uses the `zip` function and a `range`.

Note: `zip` and `enumerate` are also built-in Python functions, but their behavior is slightly different than the versions provided in `utils.py`. The behavior of the built-in variants will be described later in the course.

To check your implementation, run the following command assuming you are in the project directory once you have implemented the above functions:

When testing `utils`, please follow these steps very carefully:

Navigate to the directory containing the project: maps. Open the file named 'implementation_tester.py'. Comment out lines 2 and 3 as shown. This is to avoid getting any errors relating to implementations you have not started yet.

```
2  from abstractions import *
3  import abstractions
4  #from recommend import *
5  #import recommend
```

THESE ARE THE ONLY LINES YOU SHOULD COMMENT OR TOUCH.

To check your implementation, run the following command assuming you are in the project directory once you have implemented the above functions:

```
python3 -m unittest implementation_tester.TestProblemOne
```

Once you have passed these utils tests, go back to implementation_tester.py and UNCOMMENT these lines, as well as 4 and 5.

Problem 2

To ensure that your code passes the provided tests for abstraction, comment out these line of code in abstraction_tester.py as well as implementation_tester.py. You only have to comment out the imports in implementation_tester:

Once you start implementing Supervised and Unsupervised Learning in the later stages of the assignment, you **should uncomment import and setup statements pertaining to recommend.py**

Please note, the screenshots below show the lines that are commented that you should uncomment.

```
1  import unittest
2  import test_functions as test
3  #import recommend
4  from abstractions import *
5  import abstractions
6  #old_sample = recommend.sample
7  #test.swap_implementations(recommend)
8  #from recommend import *
```

```
369  if __name__ == '__main__':
370      unittest.main()
371      #recommend.sample = old_sample
372      #test.restore_implementations(recommend)
373      test.restore_implementations(abstractions)
374
```

Phase 1: Mean

Implement the mean function in `utils` which takes in a sequence of numbers, `s`, and returns the arithmetic mean, or average, of that sequence. The sequence cannot be empty: add an `assert` statement to ensure that empty sequences are not allowed.

Phase 2: Data Abstraction

All changes in this phase will be made to `abstractions.py`.

In this phase, we will implement data abstractions to represent a restaurant and its relevant features (such as name, location, and reviews).

Complete the implementations of the constructor and selectors for the *restaurant* data abstraction in `abstractions.py`. Two of the data abstractions have already been completed for you: the *review* data abstraction and the *user* data abstraction. Make sure that you understand how they work.

You can use any implementation you choose, but the constructor and selectors must be defined together such that the restaurant selectors return the correct field from the constructed restaurant.

- `make_restaurant`: return a restaurant constructed from five arguments:

- `name` (a string)
- `location` (a list containing latitude and longitude)
- `categories` (a list of strings)
- `price` (a number)
- `reviews` (a list of review data abstractions created by `make_review`)
- `restaurant_name`: return the name of a restaurant
- `restaurant_location`: return the location of a restaurant
- `restaurant_categories`: return the categories of a restaurant
- `restaurant_price`: return the price of a restaurant
- `restaurant_ratings`: return a list of ratings (numbers)

You must also implement the following two ratings without breaking the abstraction barrier. This means that you may not operate on the restaurant as if you know how it has been implemented. You must use the methods you have previously defined.

- `restaurant_num_ratings`: return the number of ratings restaurant
- `restaurant_mean_rating`: return the average rating of a restaurant

Test your code using:

```
python3 -m unittest implementation_tester.TestProblemTwo
```

When you finish, you should be able to generate a visualization of all restaurants rated by a user. Use `-u` to select a user from the `users` directory. You can even create your own by copying one of the `.dat` files!

```
python3 recommend.py -u one_cluster
```

Omitting the `-u` argument will default to user `test_user`.

You may have to refresh your browser to update the visualization.

Phase 2: Unsupervised Learning

All changes in this phase will be made to `recommend.py`.

Restaurants tend to appear in *clusters* (e.g. Southside restaurants, Downtown Berkeley restaurants, Gourmet Ghetto, etc.). In this phase, we will devise a way to group together restaurants that are close to each other into these clusters.

To do so, you will be implementing the **k-means algorithm**, a method for grouping data points into clusters by determining their center positions (which are called *centroids*).

K-means is considered an *unsupervised* learning method because the algorithm is not told what the correct clusters are; it must infer the clusters from the data alone.

The k-means algorithm begins by choosing k centroids at random. Then, it alternates between the following two steps:

1. **Update clusters.** Group the restaurants into clusters, where each cluster contains all restaurants that are closest to the same centroid. *In this step, centroid positions remain the same, but which cluster each restaurant belongs to can change.*
2. **Update centroids.** Compute a new centroid (average position) for each new cluster. *In this step, restaurant clusters remain the same, but the centroid positions can move.*

These steps are repeated to update the centroids until either an optimal list of centroids is found, or the centroid locations no longer change significantly each time (based on a maximum update threshold).

This visualization is a good way to understand how the algorithm works.

Glossary

As you complete the remaining questions, you will encounter the following terminology. Be sure to refer back here if you're ever confused about what a question is asking.

- **location:** A pair containing latitude and longitude. Note that this is not a data abstraction, so we can assume its implementation is a two-element list.
- **centroid:** A location (see above) that represents the center of a cluster.
- **restaurant:** A restaurant data abstraction, as defined in `abstractions.py`.
- **cluster:** A list of restaurants grouped around a centroid.
- **user:** A user data abstraction, as defined in `abstractions.py`.
- **review:** A review data abstraction, as defined in `abstractions.py`.
- **feature function:** A single argument function that takes a restaurant and returns a

- **feature function.** A single-argument function that takes a restaurant and returns a particular feature as a number, such as its mean rating or price.

Problem 3

Implement `find_closest`, which takes a `location` and a sequence of centroids (locations). It returns the element of `centroids` closest to `location`.

You should use the `distance` function from `utils.py` to measure distance between locations. The `distance` function calculates the Euclidean distance between two locations. It has been imported for you.

If two centroids are equally close, return the one that occurs first in the sequence of `centroids`.

Hint: Use the `min` function to find the centroid with the minimum distance to `location`.

Complete this google form before solving to understand how the implementation should work:

```
python3 -m unittest abstraction_tester.TestProblemTwoAbstraction
```

Problem 3 Guidance Quiz

*Required

Which of the following types of values can be passed as an argument to distance? * 1 point

- ☐ number; e.g. 1
- ☐ restaurant; e.g. `make_restaurant('A', [1, 1], ['Food'], 1, [])`
- ☐ pair; e.g. `[1, 1]`
- ☐ string of a pair; e.g. `'[1, 1]'`

Consider the list `l = [[4, 1], [-3, 2], [5, 0]]`. Which of the choices below for `fn` would make `min(l, key=fn)` evaluate to `[4, 1]`? * 1 point

- ☐ `lambda x, y: pow(-x, y)`
- ☐ `lambda x, y: abs(x - y)`
- ☐ `lambda x: abs(x[0] - x[1])`
- ☐ `sum`

SUBMIT

Google Forms

This form was created inside UC San Diego.



Test your code using:

```
python3 -m unittest implementation_tester.TestProblemThree
```

Problem 4

Now, implement a helper function for the first step in the loop of the k-means algorithm, `group_by_centroid`, which takes a sequence of restaurants and a sequence of centroids (locations) and returns a list of clusters. Each cluster of the result is a list of restaurants that are closer to a specific centroid in `centroids` than any other centroid. The order of the list of clusters returned does not matter.

If a restaurant is equally close to two centroids, it is associated with the centroid that appears first in the sequence of `centroids`.

Hint: Use the provided `group_by_first` function to group together all values for the same key in a list of `[key, value]` pairs. You can look at the doctests to see how to use it.

Complete this [google form](#) before solving to understand how the implementation should work:

Problem 4 Guidance Quiz

*Required

If centroids is $[[-1, 1], [5, -1], [1, 10], [-1, -10]]$, to which centroid will $[6, 0]$ be associated? *

1 point

- ☐ $[-1, 1]$
- ☐ $[5, -1]$
- ☐ $[1, 10]$
- ☐ $[-1, -10]$

If centroids is $[[1, 1], [1, -1], [-1, 1], [-1, -1]]$, to which centroid will $[0, 0]$ be associated? *

1 point

- ☐ $[-1, 1]$
- ☐ $[1, 1]$
- ☐ $[-1, -1]$
- ☐ $[1, -1]$

SUBMIT

Google Forms

This form was created inside UC San Diego.



Test your code using:

```
python3 -m unittest implementation_tester.TestProblemFour
```

Problem 5

Implement `find_centroid`, which finds the centroid of a `cluster` (a list of restaurants) based on the locations of the restaurants. The centroid latitude is computed by averaging the latitudes of the restaurant locations. The centroid longitude is computed by averaging the longitudes.

Hint: Use the `mean` function from `utils.py` to compute the average value of a sequence of numbers.

Be sure not violate abstraction barriers! Test your implementation before moving on:

Test your code using:

```
python3 -m unittest implementation_tester.TestProblemFive
```

```
python3 -m unittest abstraction_tester.TestProblemFiveAbstraction
```

Problem 6

Finally, implement the full `k_means` algorithm. We've already filled out the first step of the algorithm, which was to randomly initialize a list of `centroids`. The rest of the algorithm consists of iteratively updating `centroids` by grouping the restaurants into clusters based on the current list of centroids and recomputing the new centroid for each cluster.

To complete the implementation, do the following in the body of the `while` statement:

1. Group restaurants into clusters, where each cluster contains all restaurants closest to a particular centroid in `centroids`.
2. Update `centroids` to contain the true centroid (i.e. average location) for each cluster.

Hint: Use the `group_by_centroid` and `find_centroid` helper functions.

These two steps repeat until an update doesn't change the list of centroids *or* after `max_updates` iterations.

Modify or turn in any other files to complete the project. Please submit the project on Gradescope by the deadline!

`find_centroid`, which finds the centroid of a cluster (a list of restaurants) based on the locations of the restaurants. The centroid latitude is computed by averaging the latitudes of the restaurant locations. The centroid longitude is computed by averaging the longitudes.

Complete this google form before solving to understand how the implementation should work:

```
python3 -m unittest abstraction_tester.TestProblemFourAbstraction
```

Problem 6 Guidance Quiz

*Required

What is the first step of the iterative portion of the k-means algorithm? *

1 point

- ☐ randomly initialize k centroids
- ☐ create a cluster for each centroid consisting of all elements closest to that centroid
- ☐ find the centroid (average position) of each cluster.

Consider the lists `xs = [6, 1, 4]` and `ys = [2, 6, 2]`. Which of the choices below for `EXPR` would produce the following output? *

1 point

```
>>> for x, y in EXPR:
...     print(x + y)
8
7
6
```

≡ **DSC20 Fall 2019** [Syllabus](#) [Staff Hours](#) [Grading](#) [Assignments](#) ▼ [Links and](#)

- ☐ `(xs, ys)`
- ☐ `zip([xs, ys])`
- ☐ `zip(xs, ys)`

Test your code using:

```
python3 -m unittest  
abstraction_tester.TestProblemSixAbstractionAndImplementation
```

Your visualization can indicate which restaurants are close to each other (e.g. Southside restaurants, Northside restaurants). Dots that have the same color on your map belong to the same cluster of restaurants. You can get more fine-grained groupings by increasing the number of clusters with the `-k` option.

```
python3 recommend.py -k 2 python3 recommend.py -u likes_everything -k 3
```

Congratulations! You've now implemented an *unsupervised learning algorithm*.


Phase 3: Supervised Learning

All changes in this phase will be made to `recommend.py`.

In this phase, you will predict what rating a user would give for a restaurant. You will implement a *supervised* learning algorithm that attempts to generalize from examples for which the correct rating is known, which are all of the restaurants that the user has already rated. By analyzing a user's past ratings, we can then try to predict what rating the user might give to a new restaurant. When you complete this phase, your visualization will include all restaurants, not just the restaurants that were rated by a user.



To predict ratings, you will implement **simple least-squares linear regression**, a widely used statistical method that approximates a relationship between some input feature (such as price) and an output value (the rating) with a line. The algorithm takes a sequence of input-output pairs and computes the slope and intercept of the line that minimizes the mean of the squared difference between the line and the outputs.



Problem 7

First, let's use least-squares linear regression to write a function `find_predictor` that computes a predictor function for a user based on their existing restaurant ratings. A predictor function predicts a restaurant's rating based on a given feature, such as price or location. `find_predictor` takes in a user, a list of restaurants that have been reviewed by the user, and a feature function called `feature_fn` and returns two values: a predictor function and an `r_squared` value.

The predictor function is represented as the line $y = a + bx$, where y is the predicted rating for a restaurant given x , a feature of the restaurant. The `r_squared` value measures how accurately this line describes the original data.

To compute a , b , and `r_squared`, start by calculating the sums of squares S_{xx} , S_{yy} , and S_{xy} of the existing data (i.e. the feature values and corresponding user ratings for each restaurant in `restaurants`).

- $S_{xx} = \sum_i (x_i - \text{mean}(x))^2$
- $S_{yy} = \sum_i (y_i - \text{mean}(y))^2$
- $S_{xy} = \sum_i (x_i - \text{mean}(x)) (y_i - \text{mean}(y))$

After calculating the sums of squares, the regression coefficients (a and b) and `r_squared` are defined as follows:

- $b = S_{xy} / S_{xx}$
- $a = \text{mean}(y) - b * \text{mean}(x)$
- $R^2 = S_{xy}^2 / (S_{xx} S_{yy})$

Hint: The `mean` and `zip` functions can be helpful here.

Complete this google form before solving to understand how the implementation should work:

Problem 7 Guidance Quiz

*Required

What does the list `xs` represent? *

1 point

- ☐ the restaurants in restaurants
- ☐ the names of restaurants in restaurants
- ☐ the extracted values for each restaurant in restaurants
- ☐ the restaurants reviewed by user

What does the list `ys` represent? *

1 point

- ☐ the ratings for the restaurants reviewed by user
- ☐ the ratings for the restaurants in restaurants
- ☐ the names for the restaurants reviewed by user
- ☐ the names for the restaurants in restaurants

SUBMIT

Test your code using:

```
python3 -m unittest  
abstraction_tester.TestProblemSevenAbstractionAndImplementation
```

Do not worry that you are not testing implementation here, it is included in the abstraction testing.

Problem 8

Complete this google form before solving to understand how the implementation should work:

In `best_predictor`, what does the variable `reviewed` represent? *

1 point

- ☐ a list of restaurants reviewed by the user
- ☐ a list of all possible restaurants
- ☐ a list of ratings for restaurants reviewed by the user

Given a user, a list of restaurants, and a feature function, what does `find_predictor` from Problem 7 return? *

1 point

- ☐ a predictor function
- ☐ a predictor function, and its `r_squared` value
- ☐ an `r_squared` value
- ☐ a restaurant

After getting a list of `[predictor, r_squared]` pairs, which predictor should we select? *

1 point

- ☐ the first predictor in the list
- ☐ an arbitrary predictor
- ☐ the predictor with the lowest `r_squared`

Now we need a way to decide which feature is the best predictor of restaurant ratings. This can differ from user to user; for example, some users may base ratings mostly on location, while others may base them more on price.

Implement `best_predictor`, which takes a user, a list of restaurants, and a sequence of `feature_fns`. It computes a predictor function for each feature function and returns the predictor that has the highest `r_squared` value. All predictors are learned from the subset of restaurants reviewed by the user (called `reviewed` in the starter implementation).

Hint: The `max` function can also take a key argument, just like `min`.

Test your code using:

```
python3 -m unittest implementation_tester.TestProblemEight
```

```
python3 -m unittest abstraction_tester.TestProblemEightAbstraction
```

Problem 9

Complete this [google form](#) before solving to understand how the implementation should work:

rate_all returns a dictionary. What are the keys of this dictionary? *

1 point

- ☐ restaurants
- ☐ restaurant names
- ☐ restaurant ratings

What are the values of the returned dictionary? *

1 point

- ☐ numbers - user ratings only
- ☐ numbers - predicted ratings only
- ☐ numbers - a mix of user ratings and predicted ratings
- ☐ numbers - mean restaurant ratings
- ☐ lists - list of all restaurant ratings

In rate_all, what does the variable reviewed represent? *

1 point

- ☐ a list of restaurants reviewed by the user
- ☐ a list of all possible restaurants
- ☐ a list of ratings for restaurants reviewed by the user

Now that we are able to find an optimal predictor function for a given user, we can compile a full collection of restaurant ratings for the user, including ratings for restaurants they haven't actually rated yet!

Implement `rate_all`, which takes a user, a list of restaurants, and a sequence of `feature_fns`. It returns a dictionary where the keys are the names of each restaurant in `restaurants` and the values are the corresponding ratings (numbers).

If a restaurant has already been rated by the user, `rate_all` will assign the restaurant the user's rating. Otherwise, `rate_all` will assign the restaurant the rating computed by the best predictor for the user. The best predictor is chosen using a list of `feature_fns`.

Hint: `user_rating`, implemented in `abstractions.py`, returns a user's rating for a restaurant given the user and `restaurant_name`.

Be sure not violate abstraction barriers! Test your implementation before moving on:

Test your code using:

```
python3 -m unittest implementation_tester.TestProblemNine
```

```
python3 -m unittest abstraction_tester.TestProblemNineAbstraction
```

In your visualization, you can now predict what rating a user would give a restaurant, even if they haven't rated the restaurant before. To do this, add the `-p` option:

```
python3 recommend.py -u likes_southside -k 5 -p
```

If you hover over each dot (a restaurant) in the visualization, you'll see a rating in parentheses next to the restaurant name.

Problem 10

Complete this [google form](#) before solving to understand how the implementation should work:

Problem 10 Guidance Quiz

*Required

Given a restaurant, what does `restaurant_categories` in `abstractions.py` return? *

1 point

- ☐ a list of strings (categories)
- ☐ a single string (category)
- ☐ a single number (rating)
- ☐ a list of numbers (ratings)

When does a restaurant match a search query? *

1 point

- ☐ if the query string is a substring of the restaurant's name
- ☐ if the query string is mentioned in the restaurant's reviews
- ☐ if the query string is one of the restaurant's categories
- ☐ if the query string is equal to the restaurant's categories

As a final addition to our visualization, let's add the ability to focus the visualization on a particular restaurant category by implementing `search`. The `search` function takes a category query and a sequence of restaurants. It returns all restaurants that have query as a category.

Be sure not violate abstraction barriers! Test your implementation:

Test your code using:

`author?=<name>withtest_implementation_tester TestBook1.pyTest`

```
python3 -m unittest implementation_tester.TestProblemTen
```

Congratulations, you've completed the project! The `-q` option allows you to filter based on a category. For example, the following command visualizes all sandwich restaurants and their predicted ratings for the user who `likes_expensive` restaurants:

```
python3 recommend.py -u likes_expensive -k 2 -p -q Sandwiches
```

Predicting your own ratings

Once you're done, you can use your project to predict your own ratings too! Here's how:

1. In the `users` directory, you'll see a couple of `.dat` files. Copy one of them and rename the new file to `yourname.dat` (for example, `john.dat`).

1. In the new file (e.g. `john.dat`), you'll see something like the following:

```
1. make_user(      'John DoeNero',      # name      [
    # reviews      make_review('Jasmine Thai', 4.0),      ...
    ]
```

2. Replace the second line with your name (as a string).

2. Replace the existing reviews with reviews of your own! You can get a list of Berkeley restaurants with the following command:

1. `python3 recommend.py -r`

2. Rate a couple of your favorite (or least favorite) restaurants.

3. Use `recommend.py` to predict ratings for you:

1. `python3 recommend.py -u john -k 2 -p -q Sandwiches`

2. (Replace `john` with your name.) Play around with the number of clusters (the `-k` option) and try different queries (with the `-q` option)!

How accurate is your predictor?

```
python3 -m unittest abstraction_tester.TestProblemTenAbstraction
```